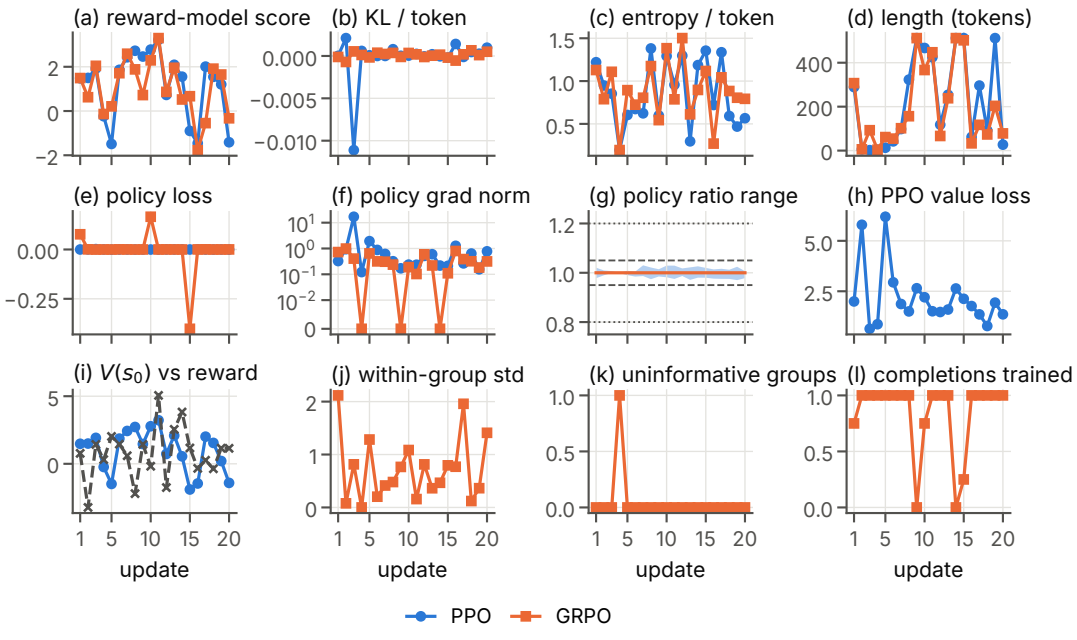
(a) reward-model score
(b) KL / token
(c) entropy / token
(d) length (tokens)
(e) policy loss
(f) policy grad norm
(g) policy ratio range
(h) PPO value loss
(i) $V(s_0)$ vs reward
(j) within-group std
(k) uninformative groups
(l) completions trained
update
PPO
GRPO