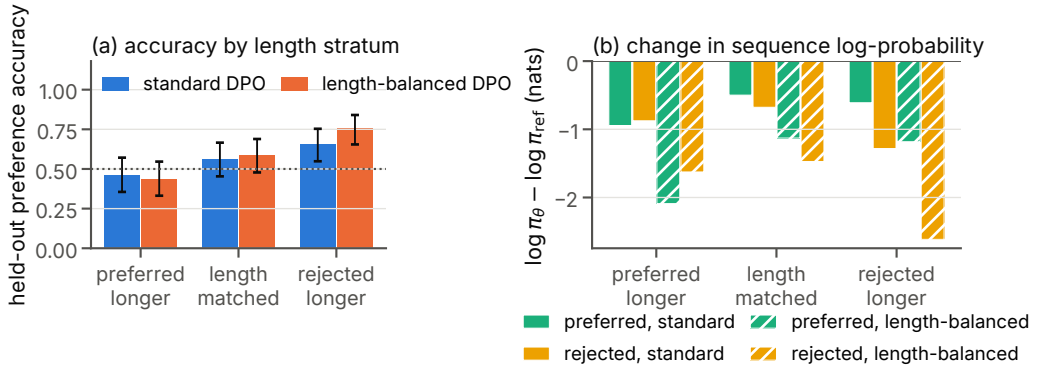
(a) accuracy by length stratum
held-out preference accuracy
standard DPO
length-balanced DPO
1.00
0.75
0.50
0.25
0.00
preferred longer
length matched
rejected longer
(b) change in sequence log-probability
$\log \pi_\theta - \log \pi_{\text{ref}}$ (nats)
0
−1
−2
preferred longer
length matched
rejected longer
preferred, standard
preferred, length-balanced
rejected, standard
rejected, length-balanced