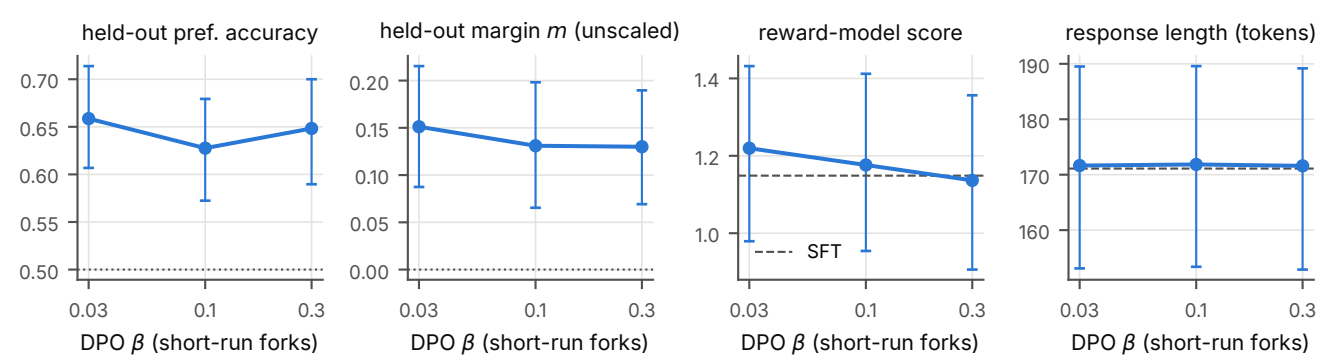
held-out pref. accuracy
0.70
0.65
0.60
0.55
0.50
0.03
0.1
0.3
DPO β (short-run forks)
held-out margin m (unscaled)
0.20
0.15
0.10
0.05
0.00
0.03
0.1
0.3
DPO β (short-run forks)
reward-model score
1.4
1.2
1.0
SFT
0.03
0.1
0.3
DPO β (short-run forks)
response length (tokens)
190
180
170
160
0.03
0.1
0.3
DPO β (short-run forks)