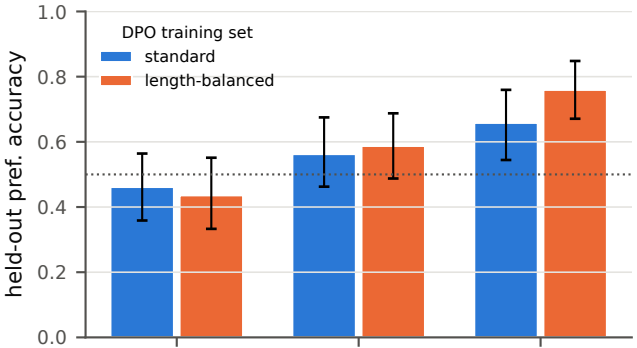
(a) accuracy by length stratum
held-out pref. accuracy
1.0
0.8
0.6
0.4
0.2
0.0
DPO training set
standard
length-balanced
preferred-longer
matched
rejected-longer

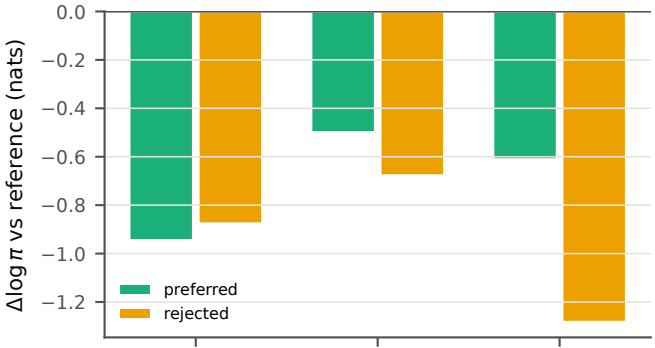
(b) standard DPO: log-prob change per response
Δlog π vs reference (nats)
0.0
−0.2
−0.4
−0.6
−0.8
−1.0
−1.2
preferred
rejected
preferred-longer
matched
rejected-longer