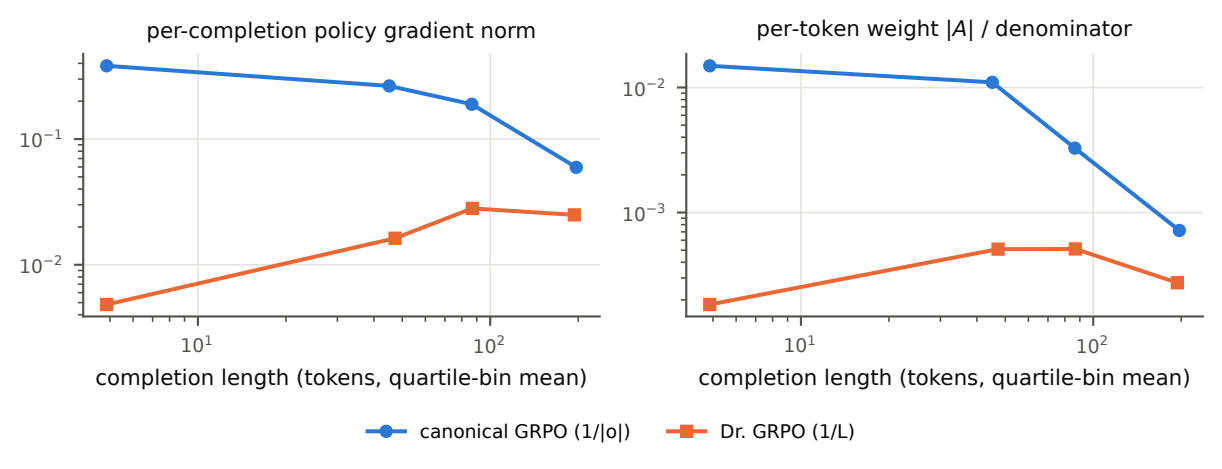
per-completion policy gradient norm
10^{-1}
10^{-2}
10^{1}
10^{2}
completion length (tokens, quartile-bin mean)
per-token weight |A| / denominator
10^{-2}
10^{-3}
10^{1}
10^{2}
completion length (tokens, quartile-bin mean)
canonical GRPO (1/|o|)
Dr. GRPO (1/L)